

```
In [15]: #Import Libraries and Load Data
import pandas as pd
import numpy as np
from sklearn.model_selection import StratifiedKFold, train_test_split
from catboost import CatBoostClassifier
from sklearn.metrics import roc_auc_score
import optuna
import warnings
import matplotlib.pyplot as plt
import seaborn as sns

warnings.filterwarnings('ignore')

# Load data
train_df = pd.read_csv("/Users/aryaman/Desktop/Spring2025/CS767/Kaggle#1/train_data.c
test_df = pd.read_csv("/Users/aryaman/Desktop/Spring2025/CS767/Kaggle#1/test_data.csv
```

```
In [16]: #Exploratory Data Analysis (EDA)
# Display basic information about the datasets
print("Train Data Info:")
print(train_df.info())
print("\nTest Data Info:")
print(test_df.info())

# Display the first few rows of the datasets
print("\nTrain Data Head:")
print(train_df.head())
print("\nTest Data Head:")
print(test_df.head())

# Check for missing values
print("\nMissing Values in Train Data:")
print(train_df.isnull().sum())
print("\nMissing Values in Test Data:")
print(test_df.isnull().sum())

# Basic statistics of the datasets
print("\nTrain Data Description:")
print(train_df.describe())
print("\nTest Data Description:")
print(test_df.describe())
```

Train Data Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29675 entries, 0 to 29674
Columns: 247 entries, game_time to ID
dtypes: bool(1), float64(30), int64(215), object(1)
memory usage: 55.7+ MB
None

Test Data Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Columns: 246 entries, game_time to ID
dtypes: float64(30), int64(215), object(1)
memory usage: 18.8+ MB
None

Train Data Head:

	game_time	game_mode	lobby_type	objectives_len	chat_len	r1_hero_id	\
0	871	22	0	4	2	110	
1	2549	22	0	17	0	114	
2	1841	22	0	8	1	100	
3	2211	22	7	11	3	32	
4	458	22	7	1	0	68	

	r1_kills	r1_deaths	r1_assists	r1_denies	...	d5_camps_stacked	\
0	2	3	11	3	...	0	
1	16	2	12	24	...	1	
2	2	11	12	2	...	0	
3	14	3	11	21	...	0	
4	3	0	0	15	...	0	

	d5_rune_pickups	d5_firstblood_claimed	d5_teamfight_participation	\
0	10	1	0.875000	
1	10	0	0.535714	
2	13	0	0.727273	
3	1	0	0.347826	
4	6	0	0.000000	

	d5_towers_killed	d5_roshans_killed	d5_obs_placed	d5_sen_placed	\
0	0	0	6	3	
1	1	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	1	1	

	radiant_win	ID
0	True	a363534a6344f1b0be1d7ba2c4047d9a
1	True	a0ba4ef0965f56d2eba69c2b9ef33353
2	True	18873e56c2142af326b4e08ca41df63a
3	True	c143931a6a8b3fb55a8ef6b9f30c6933
4	True	5a324d8b37522e9f9684493465720023

[5 rows x 247 columns]

Test Data Head:

	game_time	game_mode	lobby_type	objectives_len	chat_len	r1_hero_id	\
0	155	22	7	1	11	11	
1	1362	4	0	6	4	39	
2	2388	4	0	16	10	103	
3	2043	22	0	15	49	44	
4	840	22	7	2	27	9	

	r1_kills	r1_deaths	r1_assists	r1_denies	...	d5_creeps_stacked	\
0	0	0	0	0	...	0	
1	1	1	4	13	...	9	
2	9	8	18	3	...	0	

3	3	6	3	7	...	0
4	1	2	4	2	...	0

	d5_camps_stacked	d5_rune_pickups	d5_firstblood_claimed	\
0	0	0	0	
1	3	13	0	
2	0	6	0	
3	0	4	0	
4	0	1	0	

	d5_teamfight_participation	d5_towers_killed	d5_roshans_killed	\
0	0.000000	0	0	
1	0.222222	0	0	
2	0.686275	2	0	
3	0.945946	2	0	
4	0.375000	1	0	

	d5_obs_placed	d5_sen_placed	ID
0	0	0	a400b8f29dece5f4d266f49f1ae2e98a
1	0	0	34c81a8faede0d8f1f87dcc6ee824658
2	0	0	5feece770ca79e5e8cd8052198b3f533
3	1	0	8f56cc2468ba5c37edb79f3a7b4af6e6
4	0	0	44cdded6d3311134563f743eb77685b2

[5 rows x 246 columns]

Missing Values in Train Data:

game_time	0
game_mode	0
lobby_type	0
objectives_len	0
chat_len	0
..	
d5_roshans_killed	0
d5_obs_placed	0
d5_sen_placed	0
radiant_win	0
ID	0

Length: 247, dtype: int64

Missing Values in Test Data:

game_time	0
game_mode	0
lobby_type	0
objectives_len	0
chat_len	0
..	
d5_towers_killed	0
d5_roshans_killed	0
d5_obs_placed	0
d5_sen_placed	0
ID	0

Length: 246, dtype: int64

Train Data Description:

	game_time	game_mode	lobby_type	objectives_len	chat_len	\
count	29675.000000	29675.000000	29675.000000	29675.000000	29675.000000	
mean	1145.695333	19.592957	4.772738	6.531997	7.351879	
std	768.974419	6.297211	3.260444	6.500606	13.561608	
min	0.000000	2.000000	0.000000	0.000000	0.000000	
25%	518.000000	22.000000	0.000000	1.000000	0.000000	
50%	1043.000000	22.000000	7.000000	4.000000	3.000000	
75%	1654.000000	22.000000	7.000000	10.000000	9.000000	
max	4742.000000	23.000000	7.000000	43.000000	291.000000	

r1_hero_id	r1_kills	r1_deaths	r1_assists	r1_denies	\
------------	----------	-----------	------------	-----------	---

count	29675.000000	29675.000000	29675.000000	29675.000000	29675.000000
mean	51.103960	3.155619	3.256917	4.666992	6.335771
std	34.590915	3.744580	3.279818	5.225634	8.273026
min	1.000000	0.000000	0.000000	0.000000	0.000000
25%	20.000000	0.000000	1.000000	1.000000	1.000000
50%	44.000000	2.000000	2.000000	3.000000	3.000000
75%	81.000000	5.000000	5.000000	7.000000	9.000000
max	120.000000	32.000000	27.000000	40.000000	84.000000

	...	d5_stuns	d5_creeps_stacked	d5_camps_stacked	\
count	...	29675.000000	29675.000000	29675.000000	
mean	...	11.722391	1.038585	0.341769	
std	...	20.658919	3.556953	0.962948	
min	...	-6.191284	0.000000	0.000000	
25%	...	0.000000	0.000000	0.000000	
50%	...	1.366956	0.000000	0.000000	
75%	...	15.863060	0.000000	0.000000	
max	...	277.618070	132.000000	29.000000	

	d5_rune_pickups	d5_firstblood_claimed	d5_teamfight_participation	\
count	29675.000000	29675.000000	29675.000000	
mean	4.694457	0.088795	0.414476	
std	4.645499	0.284453	0.267860	
min	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.230769	
50%	3.000000	0.000000	0.440000	
75%	7.000000	0.000000	0.600000	
max	57.000000	1.000000	2.000000	

	d5_towers_killed	d5_roshans_killed	d5_obs_placed	d5_sen_placed
count	29675.000000	29675.000000	29675.000000	29675.000000
mean	0.302072	0.024971	1.276765	0.789250
std	0.740563	0.170487	2.602252	2.448526
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	1.000000	0.000000
max	9.000000	4.000000	26.000000	47.000000

[8 rows x 245 columns]

Test Data Description:

	game_time	game_mode	lobby_type	objectives_len	chat_len	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	1147.232600	19.560500	4.771200	6.503700	7.29880	
std	761.973655	6.328218	3.261152	6.467099	12.77006	
min	0.000000	2.000000	0.000000	0.000000	0.00000	
25%	528.000000	22.000000	0.000000	1.000000	0.00000	
50%	1048.500000	22.000000	7.000000	4.000000	3.00000	
75%	1660.000000	22.000000	7.000000	10.000000	9.00000	
max	4933.000000	23.000000	7.000000	41.000000	190.00000	

	r1_hero_id	r1_kills	r1_deaths	r1_assists	r1_denies	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	51.100700	3.124900	3.304100	4.679600	6.152700	
std	34.640793	3.663474	3.293616	5.224756	7.994321	
min	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	20.000000	0.000000	1.000000	1.000000	1.000000	
50%	44.000000	2.000000	2.000000	3.000000	3.000000	
75%	81.000000	5.000000	5.000000	7.000000	8.000000	
max	120.000000	31.000000	25.000000	34.000000	70.000000	

	...	d5_stuns	d5_creeps_stacked	d5_camps_stacked	\
count	...	10000.000000	10000.000000	10000.000000	
mean	...	11.924468	1.040300	0.347200	
std	...	20.337541	3.466393	0.966098	

min	...	-3.274992	0.000000	0.000000
25%	...	0.000000	0.000000	0.000000
50%	...	1.766235	0.000000	0.000000
75%	...	16.462201	0.000000	0.000000
max	...	235.240020	75.000000	17.000000

	d5_rune_pickups	d5_firstblood_claimed	d5_teamfight_participation	\
count	10000.000000	10000.000000	10000.000000	
mean	4.652600	0.094100	0.420367	
std	4.636536	0.291982	0.266596	
min	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.250000	
50%	3.000000	0.000000	0.444444	
75%	7.000000	0.000000	0.600000	
max	42.000000	1.000000	2.000000	


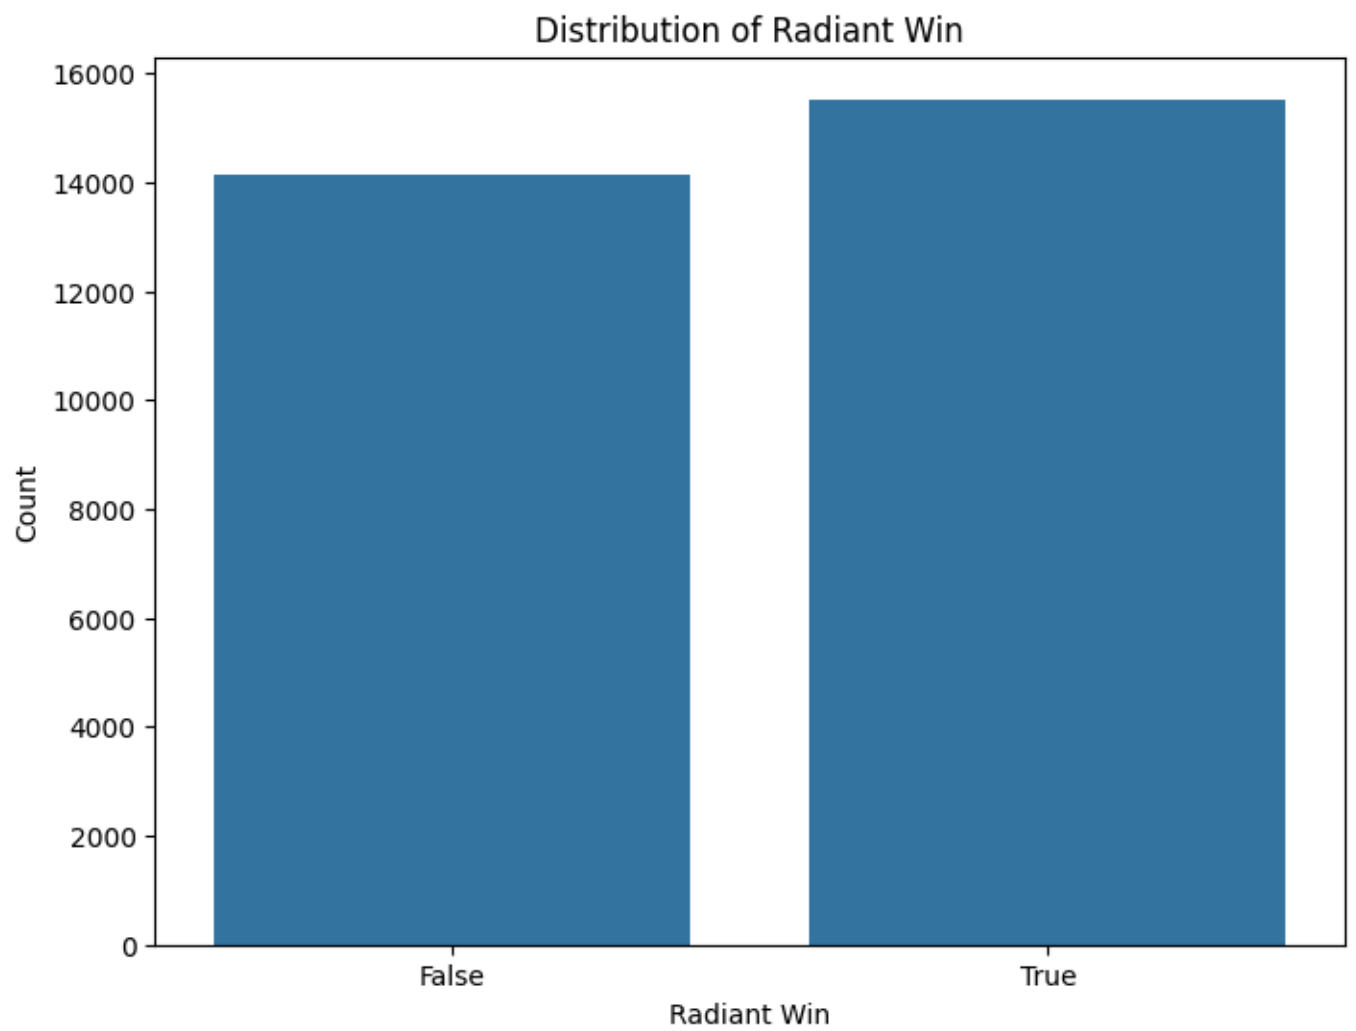
	d5_towers_killed	d5_roshans_killed	d5_obs_placed	d5_sen_placed
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.293100	0.022800	1.247100	0.765600
std	0.707985	0.170538	2.583546	2.406332
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	1.000000	0.000000
max	7.000000	5.000000	25.000000	31.000000

[8 rows x 245 columns]

```
In [17]: #Distribution of Target Variable and Game Time

# Plot the distribution of the target variable 'radiant_win'
plt.figure(figsize=(8, 6))
sns.countplot(x='radiant_win', data=train_df)
plt.title('Distribution of Radiant Win')
plt.xlabel('Radiant Win')
plt.ylabel('Count')
plt.show()

# Plot the distribution of game time
plt.figure(figsize=(8, 6))
sns.histplot(train_df['game_time'], bins=30, kde=True)
plt.title('Distribution of Game Time')
plt.xlabel('Game Time')
plt.ylabel('Frequency')
plt.show()
```




In [19]: `#Correlation Matrix`

```

numeric_cols=train_df.select_dtypes(include=[np.number]).columns.tolist()
correlation_matrix = train_df[numeric_cols].corr()

# Plot the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()

```



```

In [23]: # Ensure 'radiant_win' is treated as a numeric column
train_df['radiant_win'] = train_df['radiant_win'].astype(float)

# Ensure all columns are numeric for correlation matrix calculation
numeric_cols = train_df.select_dtypes(include=[np.number]).columns.tolist()
correlation_matrix = train_df[numeric_cols].corr()

# Correlation with target variable 'radiant_win'
if 'radiant_win' in correlation_matrix.columns:
    correlation_with_target = correlation_matrix['radiant_win'].sort_values(ascending=False)
    print("\nCorrelation with Target Variable 'radiant_win':")
    print(correlation_with_target)

# Plot correlation with target variable 'radiant_win'
plt.figure(figsize=(10, 8))
correlation_with_target.drop('radiant_win').plot(kind='bar')
plt.title('Correlation with Target Variable "radiant_win"')
plt.xlabel('Features')

```



```
plt.ylabel('Correlation Coefficient')
plt.show()
else:
    print("The target variable 'radiant_win' is not present in the correlation matrix")
```


Correlation with Target Variable 'radiant_win':

```
radiant_win      1.000000
d3_y              0.210176
r4_y              0.209289
r1_y              0.207705
r5_y              0.207584
```

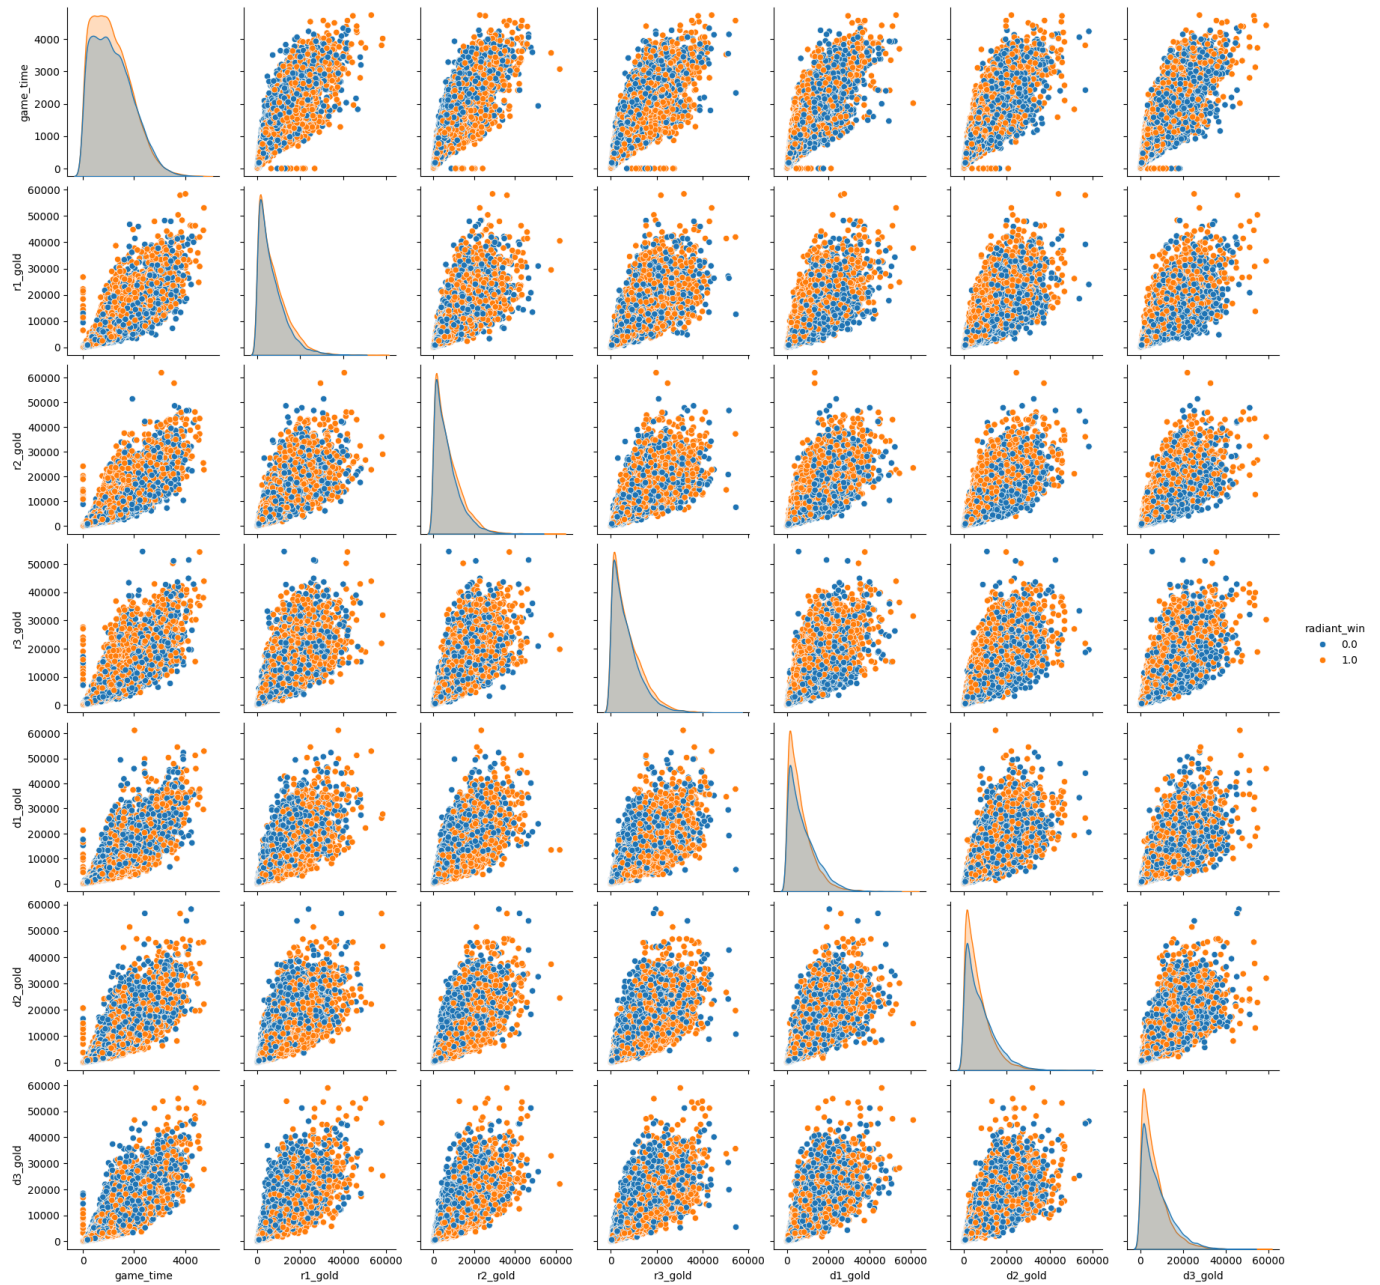
...

```
r1_deaths        -0.131264
r5_deaths        -0.133362
r3_deaths        -0.133410
d3_towers_killed -0.133541
r4_deaths        -0.135563
```

Name: radiant_win, Length: 246, dtype: float64



```
In [24]: # Pairplot for selected features
selected_features = ['game_time', 'r1_gold', 'r2_gold', 'r3_gold', 'd1_gold', 'd2_gol
sns.pairplot(train_df[selected_features], hue='radiant_win', diag_kind='kde')
plt.show()
```

```
In [25]: #Precompute Hero Win Rates
def compute_hero_win_rates(train_df):
    hero_stats = {}
    for _, row in train_df.iterrows():
        radiant_win = row['radiant_win']
        for team in ['r', 'd']:
            for pos in range(1, 6):
                hero_id = row[f'{team}{pos}_hero_id']
                win = radiant_win if team == 'r' else not radiant_win
                if hero_id not in hero_stats:
                    hero_stats[hero_id] = {'wins': 0, 'total': 0}
                hero_stats[hero_id]['wins'] += int(win)
                hero_stats[hero_id]['total'] += 1

    hero_win_rates = {hid: stats['wins']/stats['total'] for hid, stats in hero_stats.items()}
    return hero_win_rates

hero_win_rates = compute_hero_win_rates(train_df)
```

```
In [26]: #Create Enhanced Features
def create_enhanced_features(df):
    features = {}
    game_time = df['game_time']

    # Time decay factors
    early_weight = np.exp(-game_time/900)
```

```
mid_weight = np.exp(-game_time/1800)
```

```
for team in ['r', 'd']:
    # Core positions (1-3)
    core_gold = df[[f'{team}{i}_gold' for i in range(1,4)]].sum(axis=1)
    core_xp = df[[f'{team}{i}_xp' for i in range(1,4)]].sum(axis=1)

    # Support positions (4-5)
    support_gold = df[[f'{team}{i}_gold' for i in range(4,6)]].sum(axis=1)
    support_xp = df[[f'{team}{i}_xp' for i in range(4,6)]].sum(axis=1)

    # Ultra Early Game Metrics
    features[f'{team}_core_early_gold'] = (core_gold / game_time) * early_weight
    features[f'{team}_core_early_xp'] = (core_xp / game_time) * early_weight * 2.
    features[f'{team}_support_early_gold'] = (support_gold / game_time) * early_weight
    features[f'{team}_support_early_xp'] = (support_xp / game_time) * early_weight

    # Mid Game Metrics
    total_gold = core_gold + support_gold
    total_xp = core_xp + support_xp
    features[f'{team}_mid_gold'] = (total_gold / game_time) * mid_weight * 1.2
    features[f'{team}_mid_xp'] = (total_xp / game_time) * mid_weight * 1.2

    # Hero Efficiency Metrics
    for pos in range(1,6):
        # Farm efficiency with position weight
        pos_weight = 1.2 if pos <= 3 else 0.8
        features[f'{team}{pos}_farm_eff'] = (
            df[f'{team}{pos}_gold'] * 0.7 +
            df[f'{team}{pos}_lh'] * 0.3
        ) / game_time * pos_weight

        # Hero win rate features
        features[f'{team}{pos}_hero_win_rate'] = df[f'{team}{pos}_hero_id'].map(h
        features[f'{team}{pos}_hero_id'] = df[f'{team}{pos}_hero_id']

    # Team Composition Strength
    features[f'{team}_avg_win_rate'] = np.mean([
        features[f'{team}{pos}_hero_win_rate'] for pos in range(1,6)
    ], axis=0)

    # Objective Control
    features[f'{team}_objective_score'] = (
        df[[f'{team}{i}_towers_killed' for i in range(1,6)]].sum(axis=1) * 0.6 +
        df[[f'{team}{i}_roshans_killed' for i in range(1,6)]].sum(axis=1) * 0.4
    )

    # Advantage Calculations
    for metric in ['core_early_gold', 'core_early_xp', 'support_early_gold',
        'support_early_xp', 'mid_gold', 'mid_xp', 'avg_win_rate']:
        features[f'{metric}_diff'] = features[f'r_{metric}'] - features[f'd_{metric}']
        features[f'{metric}_ratio'] = (features[f'r_{metric}'] + 1e-6) / (features[f'r_{metric}'] + features[f'd_{metric}'] + 1e-6)

    # Game State Features
    features['game_time'] = game_time
    features['game_time_sq'] = game_time ** 2
    features['game_mode'] = df['game_mode']

    return pd.DataFrame(features)

# Prepare data
X_train = create_enhanced_features(train_df)
X_test = create_enhanced_features(test_df)
y_train = train_df['radiant_win']
```

```
In [27]: #Define Categorical Features  
cat_features = [col for col in X_train.columns if 'hero_id' in col or 'game_mode' in
```

```
In [28]: #Define Objective Function for Optuna  
def objective(trial):  
    params = {  
        'iterations': trial.suggest_int('iterations', 1000, 10000),  
        'learning_rate': trial.suggest_loguniform('learning_rate', 0.001, 0.1),  
        'depth': trial.suggest_int('depth', 4, 10),  
        'l2_leaf_reg': trial.suggest_loguniform('l2_leaf_reg', 1, 10),  
        'min_data_in_leaf': trial.suggest_int('min_data_in_leaf', 20, 100),  
        'random_strength': trial.suggest_uniform('random_strength', 0.1, 1.0),  
        'bagging_temperature': trial.suggest_uniform('bagging_temperature', 0.1, 1.0),  
        'grow_policy': trial.suggest_categorical('grow_policy', ['SymmetricTree', 'De  
        'random_seed': 42,  
        'verbose': False  
    }  
    model = CatBoostClassifier(**params)  
  
    X_train_split, X_val_split, y_train_split, y_val_split = train_test_split(X_train  
  
    model.fit(  
        X_train_split,  
        y_train_split,  
        eval_set=(X_val_split, y_val_split),  
        early_stopping_rounds=200,  
        cat_features=cat_features,  
        verbose=False  
    )  
  
    val_pred = model.predict_proba(X_val_split)[:, 1]  
    auc = roc_auc_score(y_val_split, val_pred)  
  
    return auc
```

```
In [29]: #Run Optuna Optimization  
study = optuna.create_study(direction='maximize')  
study.optimize(objective, n_trials=50)  
  
# Get the best parameters  
best_params = study.best_params
```

[I 2025-02-12 19:50:30,198] A new study created in memory with name: no-name-8784bb88-f075-416d-9d8d-a9c7063aeb0e

[I 2025-02-12 19:50:44,461] Trial 0 finished with value: 0.840116819500469 and parameters: {'iterations': 6108, 'learning_rate': 0.016136899498313415, 'depth': 4, 'l2_leaf_reg': 8.446318308362173, 'min_data_in_leaf': 22, 'random_strength': 0.8666159283141688, 'bagging_temperature': 0.6906527112875646, 'grow_policy': 'Lossguide'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 19:51:04,311] Trial 1 finished with value: 0.8391678231883681 and parameters: {'iterations': 8581, 'learning_rate': 0.0070169087128931435, 'depth': 6, 'l2_leaf_reg': 1.4622884021659195, 'min_data_in_leaf': 51, 'random_strength': 0.186827261534519, 'bagging_temperature': 0.5330639518753737, 'grow_policy': 'Lossguide'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 19:55:58,528] Trial 2 finished with value: 0.838903865456193 and parameters: {'iterations': 5951, 'learning_rate': 0.0010325389691269178, 'depth': 10, 'l2_leaf_reg': 1.0235172131803476, 'min_data_in_leaf': 27, 'random_strength': 0.14672346173641235, 'bagging_temperature': 0.11257164111035663, 'grow_policy': 'SymmetricTree'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 19:56:01,823] Trial 3 finished with value: 0.8382963075726517 and parameters: {'iterations': 3835, 'learning_rate': 0.06024574635936515, 'depth': 5, 'l2_leaf_reg': 2.643255667939475, 'min_data_in_leaf': 24, 'random_strength': 0.3793038468797767, 'bagging_temperature': 0.561516469110466, 'grow_policy': 'Depthwise'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 19:59:31,531] Trial 4 finished with value: 0.8395877662872729 and parameters: {'iterations': 5054, 'learning_rate': 0.0014040764357510708, 'depth': 8, 'l2_leaf_reg': 7.263857393006822, 'min_data_in_leaf': 21, 'random_strength': 0.3512751085266774, 'bagging_temperature': 0.8172135363589489, 'grow_policy': 'Lossguide'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 20:01:40,280] Trial 5 finished with value: 0.8393631746651549 and parameters: {'iterations': 5025, 'learning_rate': 0.001440363885708944, 'depth': 8, 'l2_leaf_reg': 2.691391948291677, 'min_data_in_leaf': 31, 'random_strength': 0.11173220196351924, 'bagging_temperature': 0.34856862638493513, 'grow_policy': 'SymmetricTree'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 20:02:09,375] Trial 6 finished with value: 0.8397786805435616 and parameters: {'iterations': 9841, 'learning_rate': 0.01532091517656213, 'depth': 8, 'l2_leaf_reg': 1.8090475798680163, 'min_data_in_leaf': 86, 'random_strength': 0.9520780655340746, 'bagging_temperature': 0.45197168483650363, 'grow_policy': 'SymmetricTree'}. Best is trial 0 with value: 0.840116819500469.

[I 2025-02-12 20:02:30,533] Trial 7 finished with value: 0.8402691640708927 and parameters: {'iterations': 5638, 'learning_rate': 0.01576771117496574, 'depth': 4, 'l2_leaf_reg': 4.276858086261457, 'min_data_in_leaf': 80, 'random_strength': 0.841672733386203, 'bagging_temperature': 0.22268708821269015, 'grow_policy': 'SymmetricTree'}. Best is trial 7 with value: 0.8402691640708927.

[I 2025-02-12 20:03:55,070] Trial 8 finished with value: 0.8392202734101495 and parameters: {'iterations': 6401, 'learning_rate': 0.007106411308149588, 'depth': 10, 'l2_leaf_reg': 4.1161448857208525, 'min_data_in_leaf': 97, 'random_strength': 0.1274194692941861, 'bagging_temperature': 0.4420618145588583, 'grow_policy': 'SymmetricTree'}. Best is trial 7 with value: 0.8402691640708927.

[I 2025-02-12 20:04:06,330] Trial 9 finished with value: 0.8369534226102097 and parameters: {'iterations': 1031, 'learning_rate': 0.08561619794116097, 'depth': 9, 'l2_leaf_reg': 4.677860954568512, 'min_data_in_leaf': 21, 'random_strength': 0.27766251458879343, 'bagging_temperature': 0.6277468871047887, 'grow_policy': 'SymmetricTree'}. Best is trial 7 with value: 0.8402691640708927.

[I 2025-02-12 20:04:12,072] Trial 10 finished with value: 0.8397828902142905 and parameters: {'iterations': 2388, 'learning_rate': 0.03302969810959035, 'depth': 4, 'l2_leaf_reg': 5.628534155558901, 'min_data_in_leaf': 72, 'random_strength': 0.7177954147034526, 'bagging_temperature': 0.14692506863080623, 'grow_policy': 'Depthwise'}. Best is trial 7 with value: 0.8402691640708927.

[I 2025-02-12 20:04:27,704] Trial 11 finished with value: 0.8403758849126084 and parameters: {'iterations': 7584, 'learning_rate': 0.01784220258669064, 'depth': 4, 'l2_leaf_reg': 9.84342974139258, 'min_data_in_leaf': 55, 'random_strength': 0.8997196864300747, 'bagging_temperature': 0.9425535690730353, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:05:18,565] Trial 12 finished with value: 0.8399693672500815 and parameters: {'iterations': 7718, 'learning_rate': 0.0036131760581170416, 'depth': 6, 'l2_leaf_reg': 8.9993262299375, 'min_data_in_leaf': 57, 'random_strength': 0.753987043402723

8, 'bagging_temperature': 0.9466880229862766, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:05:26,137] Trial 13 finished with value: 0.8391564456999123 and parameters: {'iterations': 7504, 'learning_rate': 0.03167049504096631, 'depth': 5, 'l2_leaf_reg': 3.48949600862125, 'min_data_in_leaf': 70, 'random_strength': 0.6213814905674424, 'bagging_temperature': 0.27777239327178194, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:05:31,200] Trial 14 finished with value: 0.8399768763924623 and parameters: {'iterations': 3959, 'learning_rate': 0.02720874318578251, 'depth': 4, 'l2_leaf_reg': 6.719686285227397, 'min_data_in_leaf': 44, 'random_strength': 0.966216263743576, 'bagging_temperature': 0.9775810301143932, 'grow_policy': 'Depthwise'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:07:07,781] Trial 15 finished with value: 0.8402941945454956 and parameters: {'iterations': 7546, 'learning_rate': 0.0036414548805247396, 'depth': 6, 'l2_leaf_reg': 9.781470194988852, 'min_data_in_leaf': 74, 'random_strength': 0.5335892429400164, 'bagging_temperature': 0.7760087256233625, 'grow_policy': 'SymmetricTree'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:08:06,574] Trial 16 finished with value: 0.8401566407100645 and parameters: {'iterations': 9464, 'learning_rate': 0.0033341543115904307, 'depth': 6, 'l2_leaf_reg': 5.823472631406535, 'min_data_in_leaf': 41, 'random_strength': 0.4883384202844493, 'bagging_temperature': 0.8031427717828161, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:09:49,232] Trial 17 finished with value: 0.8401403709015728 and parameters: {'iterations': 7352, 'learning_rate': 0.003085713993094283, 'depth': 7, 'l2_leaf_reg': 9.730120718149262, 'min_data_in_leaf': 65, 'random_strength': 0.5450368617923435, 'bagging_temperature': 0.8321743982644865, 'grow_policy': 'SymmetricTree'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:10:26,113] Trial 18 finished with value: 0.840263930426203 and parameters: {'iterations': 8528, 'learning_rate': 0.005924990898438478, 'depth': 5, 'l2_leaf_reg': 7.460349197299446, 'min_data_in_leaf': 82, 'random_strength': 0.6779108642577275, 'bagging_temperature': 0.7106932719772387, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:10:44,543] Trial 19 finished with value: 0.8396343002150572 and parameters: {'iterations': 6918, 'learning_rate': 0.009989465005642108, 'depth': 7, 'l2_leaf_reg': 9.99981249861367, 'min_data_in_leaf': 59, 'random_strength': 0.510230589395273, 'bagging_temperature': 0.9050795037636967, 'grow_policy': 'Depthwise'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:11:57,372] Trial 20 finished with value: 0.8399232884218355 and parameters: {'iterations': 8579, 'learning_rate': 0.0021046428439238985, 'depth': 5, 'l2_leaf_reg': 5.701078434003654, 'min_data_in_leaf': 100, 'random_strength': 0.4472636240379829, 'bagging_temperature': 0.7401224824269813, 'grow_policy': 'Lossguide'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:12:12,731] Trial 21 finished with value: 0.8401911145000858 and parameters: {'iterations': 4104, 'learning_rate': 0.01805685814670291, 'depth': 4, 'l2_leaf_reg': 4.5037419233569285, 'min_data_in_leaf': 82, 'random_strength': 0.8172189298150612, 'bagging_temperature': 0.897856410430868, 'grow_policy': 'SymmetricTree'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:12:42,216] Trial 22 finished with value: 0.8402800864598101 and parameters: {'iterations': 8057, 'learning_rate': 0.010451184725495558, 'depth': 5, 'l2_leaf_reg': 1.9787620996790816, 'min_data_in_leaf': 74, 'random_strength': 0.8748289159652083, 'bagging_temperature': 0.2194597174264517, 'grow_policy': 'SymmetricTree'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:13:32,972] Trial 23 finished with value: 0.840258014132206 and parameters: {'iterations': 8229, 'learning_rate': 0.004935128849110394, 'depth': 6, 'l2_leaf_reg': 1.909008723796015, 'min_data_in_leaf': 73, 'random_strength': 0.6381858644035501, 'bagging_temperature': 0.9905094994818276, 'grow_policy': 'SymmetricTree'}. Best is trial 11 with value: 0.8403758849126084.

[I 2025-02-12 20:14:01,372] Trial 24 finished with value: 0.8405149178215386 and parameters: {'iterations': 9160, 'learning_rate': 0.010270110683315797, 'depth': 5, 'l2_leaf_reg': 2.156218322064168, 'min_data_in_leaf': 92, 'random_strength': 0.8808335395253981, 'bagging_temperature': 0.8881353073486111, 'grow_policy': 'SymmetricTree'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:14:38,486] Trial 25 finished with value: 0.8400896273030598 and parameters: {'iterations': 9368, 'learning_rate': 0.0103371804635015, 'depth': 6, 'l2_leaf_reg': 2.4276511856981817, 'min_data_in_leaf': 93, 'random_strength': 0.9945114187048847, 'bagging_temperature': 0.8579913168153553, 'grow_policy': 'SymmetricTree'}. Best is

trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:14:55,300] Trial 26 finished with value: 0.840242540747906 and parameters: {'iterations': 6865, 'learning_rate': 0.023164423353495465, 'depth': 5, 'l2_leaf_reg': 1.439812739585279, 'min_data_in_leaf': 92, 'random_strength': 0.5957369801095421, 'bagging_temperature': 0.7658082542038224, 'grow_policy': 'SymmetricTree'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:15:07,339] Trial 27 finished with value: 0.8396038085459956 and parameters: {'iterations': 9232, 'learning_rate': 0.039778352553750604, 'depth': 7, 'l2_leaf_reg': 3.585898511140938, 'min_data_in_leaf': 65, 'random_strength': 0.7728258312325611, 'bagging_temperature': 0.6511436007191285, 'grow_policy': 'SymmetricTree'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:15:44,140] Trial 28 finished with value: 0.840312853626563 and parameters: {'iterations': 9269, 'learning_rate': 0.004741802703234369, 'depth': 5, 'l2_leaf_reg': 2.9793547059230647, 'min_data_in_leaf': 51, 'random_strength': 0.9099161253450229, 'bagging_temperature': 0.9160750216237548, 'grow_policy': 'Depthwise'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:15:57,138] Trial 29 finished with value: 0.8397549015926891 and parameters: {'iterations': 9927, 'learning_rate': 0.012719804828744479, 'depth': 4, 'l2_leaf_reg': 2.9501145609990655, 'min_data_in_leaf': 52, 'random_strength': 0.8975734611132241, 'bagging_temperature': 0.9034855915614565, 'grow_policy': 'Depthwise'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:16:09,369] Trial 30 finished with value: 0.8402581279070904 and parameters: {'iterations': 8908, 'learning_rate': 0.020995118986986222, 'depth': 4, 'l2_leaf_reg': 2.377708818618802, 'min_data_in_leaf': 40, 'random_strength': 0.9142910265384434, 'bagging_temperature': 0.9961259583270871, 'grow_policy': 'Depthwise'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:16:43,334] Trial 31 finished with value: 0.840328327010863 and parameters: {'iterations': 6999, 'learning_rate': 0.004491127513762571, 'depth': 5, 'l2_leaf_reg': 8.265034524008277, 'min_data_in_leaf': 50, 'random_strength': 0.8151020264382891, 'bagging_temperature': 0.9189719837974684, 'grow_policy': 'Depthwise'}. Best is trial 24 with value: 0.8405149178215386.

[I 2025-02-12 20:17:06,953] Trial 32 finished with value: 0.8405179897434216 and parameters: {'iterations': 6683, 'learning_rate': 0.0076517313589637365, 'depth': 5, 'l2_leaf_reg': 7.703144559140397, 'min_data_in_leaf': 48, 'random_strength': 0.7737196601730773, 'bagging_temperature': 0.8747240514805957, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:17:32,979] Trial 33 finished with value: 0.840226043389645 and parameters: {'iterations': 6432, 'learning_rate': 0.007073386615817629, 'depth': 5, 'l2_leaf_reg': 8.206620485948218, 'min_data_in_leaf': 46, 'random_strength': 0.8051472519898805, 'bagging_temperature': 0.8611268165821716, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:17:50,030] Trial 34 finished with value: 0.8397795907426383 and parameters: {'iterations': 6922, 'learning_rate': 0.00839308082381201, 'depth': 4, 'l2_leaf_reg': 8.169691149730921, 'min_data_in_leaf': 34, 'random_strength': 0.7245283436365173, 'bagging_temperature': 0.945042952854767, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:18:15,379] Trial 35 finished with value: 0.8401081726092425 and parameters: {'iterations': 6054, 'learning_rate': 0.005305396067026149, 'depth': 5, 'l2_leaf_reg': 1.095481759492582, 'min_data_in_leaf': 56, 'random_strength': 0.8347153023126155, 'bagging_temperature': 0.6851321110488148, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:18:32,085] Trial 36 finished with value: 0.8401215980456206 and parameters: {'iterations': 7928, 'learning_rate': 0.012591341330481155, 'depth': 6, 'l2_leaf_reg': 6.735289203150057, 'min_data_in_leaf': 64, 'random_strength': 0.7766095914133045, 'bagging_temperature': 0.861226323502286, 'grow_policy': 'Lossguide'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:19:08,743] Trial 37 finished with value: 0.8399490015457456 and parameters: {'iterations': 4868, 'learning_rate': 0.0024800892303757016, 'depth': 4, 'l2_leaf_reg': 1.5345902718040434, 'min_data_in_leaf': 48, 'random_strength': 0.945381813598197, 'bagging_temperature': 0.6101103205568552, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:19:15,265] Trial 38 finished with value: 0.8393292697495565 and parameters: {'iterations': 5538, 'learning_rate': 0.05101563196440233, 'depth': 6, 'l2_leaf_reg': 5.135668123447489, 'min_data_in_leaf': 39, 'random_strength': 0.6772378871223135, 'bagging_temperature': 0.5353201015632082, 'grow_policy': 'Lossguide'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:20:18,307] Trial 39 finished with value: 0.840239696375792 and parameters: {'iterations': 6615, 'learning_rate': 0.0018088001000109352, 'depth': 5, 'l2_leaf_reg': 7.500338296234908, 'min_data_in_leaf': 35, 'random_strength': 0.8726668757787366, 'bagging_temperature': 0.804800017489462, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:20:39,288] Trial 40 finished with value: 0.8401457183211468 and parameters: {'iterations': 7322, 'learning_rate': 0.008160642984981068, 'depth': 4, 'l2_leaf_reg': 6.4347012080855, 'min_data_in_leaf': 55, 'random_strength': 0.9944869104685158, 'bagging_temperature': 0.45786364357642956, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:21:16,925] Trial 41 finished with value: 0.840160622831024 and parameters: {'iterations': 8941, 'learning_rate': 0.004544215343728066, 'depth': 5, 'l2_leaf_reg': 2.3238147984103152, 'min_data_in_leaf': 51, 'random_strength': 0.9200146608915686, 'bagging_temperature': 0.9385870913509222, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:21:52,636] Trial 42 finished with value: 0.8402812242086558 and parameters: {'iterations': 8202, 'learning_rate': 0.004293562354736035, 'depth': 5, 'l2_leaf_reg': 3.0982394849336874, 'min_data_in_leaf': 49, 'random_strength': 0.8578982625845509, 'bagging_temperature': 0.898450105122085, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:22:30,626] Trial 43 finished with value: 0.8396456777035131 and parameters: {'iterations': 5868, 'learning_rate': 0.005891483970324492, 'depth': 9, 'l2_leaf_reg': 3.799049019655831, 'min_data_in_leaf': 61, 'random_strength': 0.7975919198897504, 'bagging_temperature': 0.945066289802122, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:22:45,803] Trial 44 finished with value: 0.8400666447763789 and parameters: {'iterations': 8994, 'learning_rate': 0.014278129241013316, 'depth': 5, 'l2_leaf_reg': 2.7524826388084187, 'min_data_in_leaf': 53, 'random_strength': 0.7332238269910353, 'bagging_temperature': 0.8588945207624913, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:23:48,251] Trial 45 finished with value: 0.8402437922716361 and parameters: {'iterations': 9601, 'learning_rate': 0.0026567971597554283, 'depth': 4, 'l2_leaf_reg': 8.758902347239852, 'min_data_in_leaf': 44, 'random_strength': 0.9425331661703582, 'bagging_temperature': 0.9213702918398171, 'grow_policy': 'Lossguide'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:24:04,659] Trial 46 finished with value: 0.8401489040179144 and parameters: {'iterations': 5239, 'learning_rate': 0.008878473023351633, 'depth': 5, 'l2_leaf_reg': 2.0478981910969103, 'min_data_in_leaf': 59, 'random_strength': 0.8408596017345434, 'bagging_temperature': 0.8210202402382967, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:24:18,103] Trial 47 finished with value: 0.8399621994323544 and parameters: {'iterations': 8600, 'learning_rate': 0.017976252709374654, 'depth': 6, 'l2_leaf_reg': 1.615162234864446, 'min_data_in_leaf': 69, 'random_strength': 0.8930101485163273, 'bagging_temperature': 0.748303635980065, 'grow_policy': 'Lossguide'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:25:00,526] Trial 48 finished with value: 0.8400376321808165 and parameters: {'iterations': 7253, 'learning_rate': 0.00410967601698029, 'depth': 4, 'l2_leaf_reg': 3.334035198549425, 'min_data_in_leaf': 43, 'random_strength': 0.958180967394914, 'bagging_temperature': 0.9543447191307586, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

[I 2025-02-12 20:25:18,025] Trial 49 finished with value: 0.8394430446341147 and parameters: {'iterations': 4541, 'learning_rate': 0.007284043974688932, 'depth': 4, 'l2_leaf_reg': 7.910687688174968, 'min_data_in_leaf': 26, 'random_strength': 0.20518735359589313, 'bagging_temperature': 0.9957106705953265, 'grow_policy': 'Depthwise'}. Best is trial 32 with value: 0.8405179897434216.

```
In [30]: #Training Setup and Model Training
n_folds = 7
predictions = np.zeros(len(X_test))
oof_predictions = np.zeros(len(X_train))
skf = StratifiedKFold(n_splits=n_folds, shuffle=True, random_state=42)

for fold, (train_idx, val_idx) in enumerate(skf.split(X_train, y_train)):
    print(f"\nFold {fold + 1}:")

    X_fold_train = X_train.iloc[train_idx]
```



```

y_fold_train = y_train.iloc[train_idx]
X_fold_val = X_train.iloc[val_idx]
y_fold_val = y_train.iloc[val_idx]

model = CatBoostClassifier(**best_params)

model.fit(
    X_fold_train,
    y_fold_train,
    eval_set=(X_fold_val, y_fold_val),
    early_stopping_rounds=200,
    cat_features=cat_features,
    verbose=False
)

val_pred = model.predict_proba(X_fold_val)[:, 1]
oof_predictions[val_idx] = val_pred
print(f"Validation AUC: {roc_auc_score(y_fold_val, val_pred):.4f}")

predictions += model.predict_proba(X_test)[:, 1] / n_folds

print(f"\nOverall CV score: {roc_auc_score(y_train, oof_predictions):.4f}")

```

Fold 1:
Validation AUC: 0.8436

Fold 2:
Validation AUC: 0.8381

Fold 3:
Validation AUC: 0.8494

Fold 4:
Validation AUC: 0.8414

Fold 5:
Validation AUC: 0.8412

Fold 6:
Validation AUC: 0.8464

Fold 7:
Validation AUC: 0.8300

Overall CV score: 0.8415

```

In [31]: # Create submission
submission = pd.DataFrame({
    'ID': test_df['ID'],
    'radiant_win': predictions
})
submission.to_csv('dota2_kaggle.csv', index=False)
print("\nOptimized submission file created!")

```

Optimized submission file created!

```

In [3]: !jupyter nbconvert --to html Dota2_Kaggle.ipynb

```

```
Traceback (most recent call last):
  File "/Users/aryaman/Library/Python/3.9/bin/jupyter-nbconvert", line 5, in <module>
    from nbconvert.nbconvertapp import main
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/nbconvert/nbconvert
app.py", line 193, in <module>
    class NbConvertApp(JupyterApp):
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/nbconvert/nbconvert
app.py", line 252, in NbConvertApp
    Options include {get_export_names()}.
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/nbconvert/exporter
s/base.py", line 145, in get_export_names
    e = get_exporter(exporter_name)(config=config)
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/nbconvert/exporter
s/base.py", line 106, in get_exporter
    exporter = items[0].load()
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/importlib_metadata/
__init__.py", line 189, in load
    module = import_module(match.group('module'))
  File "/Library/Developer/CommandLineTools/Library/Frameworks/Python3.framework/Versi
ons/3.9/lib/python3.9/importlib/__init__.py", line 127, in import_module
    return _bootstrap._gcd_import(name[level:], package, level)
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/jupyter_contrib_nbe
xtensions/nbconvert_support/__init__.py", line 5, in <module>
    from .collapsible_headings import ExporterCollapsibleHeadings
  File "/Users/aryaman/Library/Python/3.9/lib/python/site-packages/jupyter_contrib_nbe
xtensions/nbconvert_support/collapsible_headings.py", line 6, in <module>
    from notebook.services.config import ConfigManager
ModuleNotFoundError: No module named 'notebook.services'
```